AWS Academy Cloud Foundations

Module 10: Automatic Scaling and Monitoring

aws academy

Module overview

Topics

- Elastic Load Balancing

- Amazon CloudWatch

- Amazon EC2 Auto Scaling

Activities

- Elastic Load Balancing activity

- Amazon CloudWatch activity

Lab

- Scale and Load Balance Your Architecture

Knowledge check

Module objectives

After completing this module, you should be able to:

- Indicate how to distribute traffic across Amazon Elastic Compute Cloud (Amazon EC2) instances by using Elastic Load Balancing

- Identify how Amazon CloudWatch enables you to monitor AWS resources and applications in real time

- Explain how Amazon EC2 Auto Scaling launches and releases servers in response to workload changes

- Perform scaling and load balancing tasks to improve an architecture

Module 10: Automatic Scaling and Monitoring

Section 1: Elastic Load Balancing

aws academy

Elastic Load Balancing

- Distributes incoming application or network traffic across multiple targets in a single Availability Zone or across multiple Availability Zones.

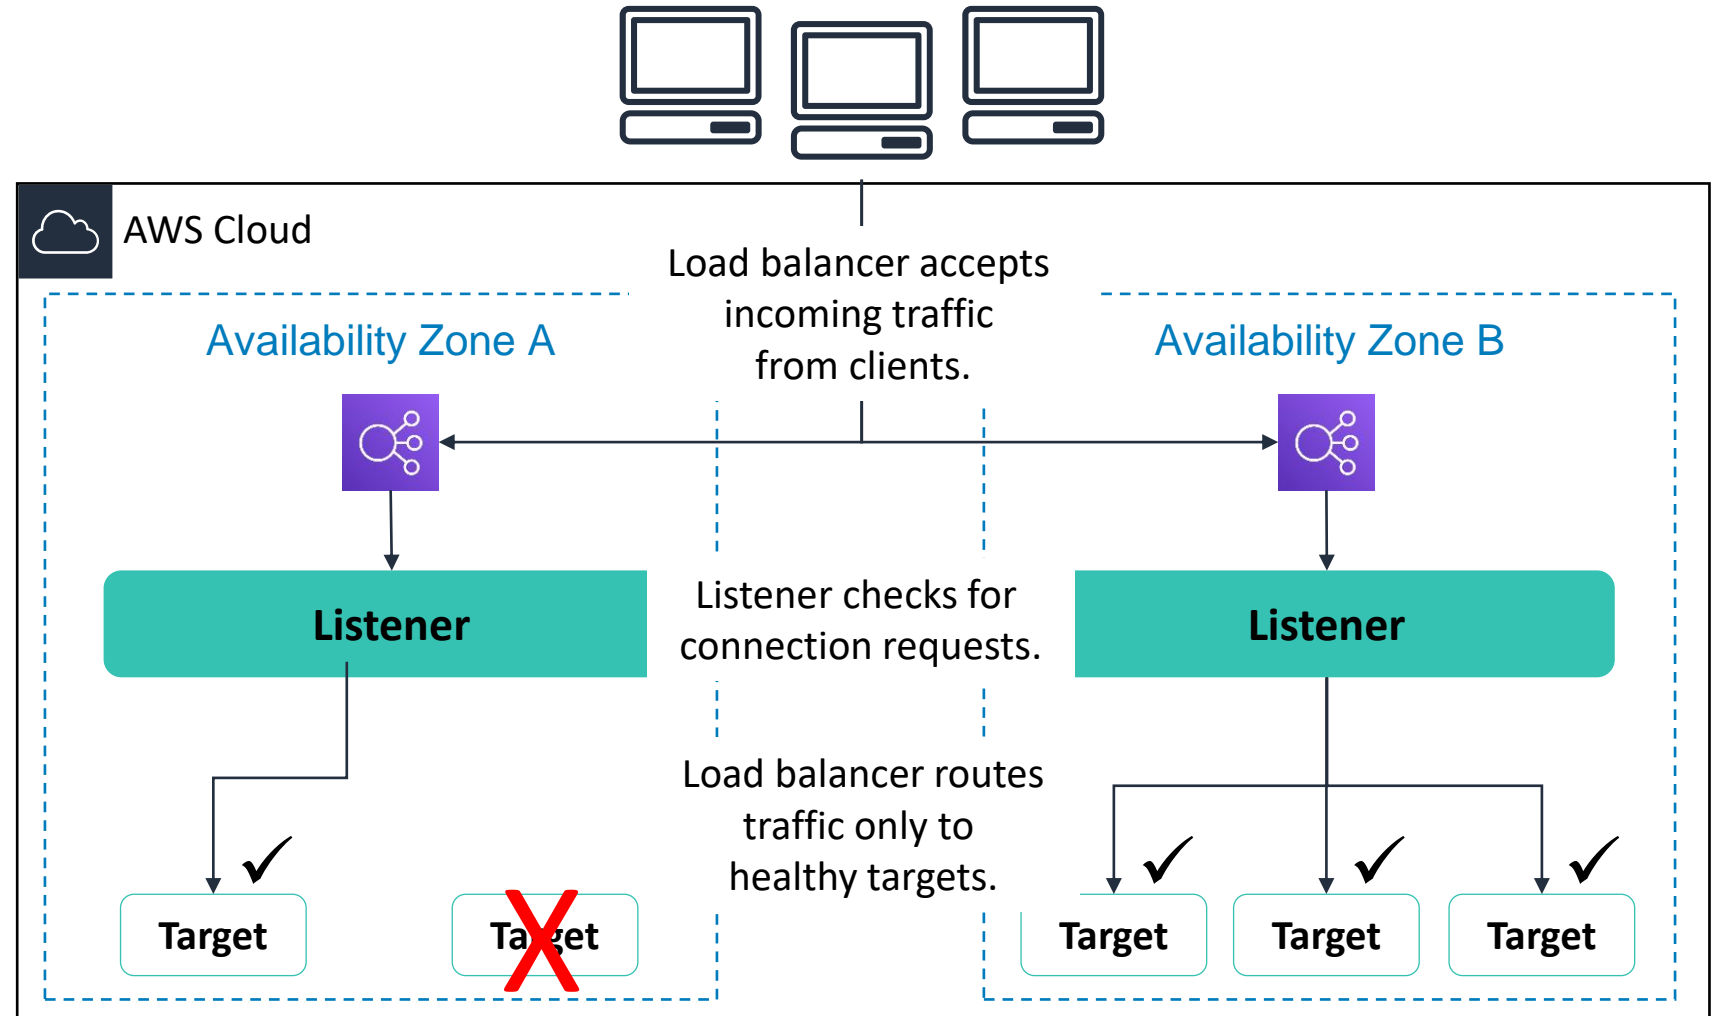- Scales your load balancer as traffic to your application changes over time.

Elastic Load Balancing

EC2 instances

Containers

IP addresses

Lambda functions

Types of load balancers

Application Load Balancer	Network Load Balancer	Classic Load Balancer (Previous Generation)
• Load balancing of HTTP and HTTPS traffic	• Load balancing of TCP, UDP, and TLS traffic where extreme performance is required	• Load balancing of HTTP, HTTPS, TCP, and SSL traffic
• Routes traffic to targets based on content of request • Provides advanced request routing targeted at the delivery of modern application architectures, including microservices and containers	• Routes traffic to targets based on IP protocol data • Can handle millions of requests per second while maintaining ultra-low latencies • Is optimized to handle sudden and volatile traffic patterns	• Load balancing across multiple EC2 instances
• Operates at the application layer (OSI model layer 7)	• Operates at the transport layer (OSI model layer 4)	• Operates at both the application and transport layers.

How Elastic Load Balancing works

- With Application Load Balancers and Network Load Balancers, you register targets in target groups, and route traffic to the target groups.

- With Classic Load Balancers, you register instances with the load balancer.

Load balancer performs health checks to monitor health of registered targets.

AWS Cloud

Load balancer accepts incoming traffic from clients.

Availability Zone A

Availability Zone B

Listener checks for connection requests.

Listener

Listener

Load balancer routes traffic only to healthy targets.

✓

Target

Target

✗

Target

✓

Target

✓

Target

✓

Elastic Load Balancing use cases

Highly available and
fault-tolerant
applications

Containerized
applications

Elasticity
and scalability

Virtual private
cloud (VPC)

Hybrid environments

Invoke Lambda
functions over HTTP(S)

Activity: Elastic Load Balancing

|	
You must support traffic to a containerized application.	Application Load Balancer
You have extremely spiky and unpredictable TCP traffic.	Network Load Balancer
You need simple load balancing with multiple protocols.	Classic Load Balancer
You need to support a static or Elastic IP address, or an IP target outside a VPC.	Network Load Balancer
You need a load balancer that can handle millions of requests per second while maintaining low latencies.	Network Load Balancer
You must support HTTPS requests.	Application Load Balancer

Load balancer monitoring

- **Amazon CloudWatch metrics** – Used to verify that the system is performing as expected and creates an alarm to initiate an action if a metric goes outside an acceptable range.

- **Access logs** – Capture detailed information about requests sent to your load balancer.

- **AWS CloudTrail logs** – Capture the who, what, when, and where of API interactions in AWS services.

Section 1 key takeaways

- Elastic Load Balancing distributes incoming application or network traffic across multiple targets in one or more Availability Zones.

- Elastic Load Balancing supports three types of load balancers:
  - Application Load Balancer
  - Network Load Balancer
  - Classic Load Balancer

- ELB offers instance health checks, security, and monitoring.

**Module 10: Automatic Scaling and Monitoring**

Section 2: Amazon CloudWatch

aws academy

Monitoring AWS resources

To use AWS efficiently, you need insight into your AWS resources:

- How do you know when you should **launch more Amazon EC2 instances**?

- Is your **application's performance or availability** being affected by a lack of sufficient capacity?

- How much of your infrastructure is actually **being used**?

13

Amazon CloudWatch

Amazon CloudWatch

- Monitors –
  - AWS resources
  - Applications that run on AWS
- Collects and tracks –
  - Standard metrics
  - Custom metrics
- Alarms –
  - Send notifications to an Amazon SNS topic
  - Perform Amazon EC2 Auto Scaling or Amazon EC2 actions
- Events –
  - Define rules to match changes in AWS environment and route these events to one or more target functions or streams for processing

CloudWatch alarms

- Create alarms based on –
  - Static threshold
  - Anomaly detection
  - Metric math expression
- Specify –
  - Namespace
  - Metric
  - Statistic
  - Period
  - Conditions
  - Additional configuration
  - Actions

Statistic

🔍 Average ✕

Period

5 minutes ▼

Conditions

Threshold type

⦿ Static
Use a value as a threshold

○ Anomaly detection
Use a band as a threshold

Whenever CPUUtilization is...
Define the alarm condition

⦿ Greater
\> threshold

○ Greater/Equal
\>= threshold

○ Lower/Equal
<= threshold

○ Lower
< threshold

than...
Define the threshold value

100 ⬍

Must be a number

▶ Additional configuration

Activity: Amazon CloudWatch

Amazon EC2

| If average CPU utilization is > 60% for 5 minutes… | Correct! |

Amazon RDS

| If the number of simultaneous connections is > 10 for 1 minute… | Correct! |

Amazon S3

| If the maximum bucket size in bytes is around 3 for 1 day… | Incorrect. *Around* is not a threshold option. You must specify a threshold of >, >=, <=, or <. |

Elastic Load Balancing

| If the number of healthy hosts is < 5 for 10 minutes… | Correct! |

Amazon Elastic Block Store

| If the volume of read operations is > 1,000 for 10 seconds… | Incorrect. You must specify a statistic (for example, *average volume*). |

Section 2 key takeaways

- Amazon CloudWatch helps you monitor your AWS resources—and the applications that you run on AWS—in real time.

- CloudWatch enables you to –
  - Collect and track standard and custom metrics.
  - Set alarms to automatically send notifications to SNS topics, or perform Amazon EC2 Auto Scaling or Amazon EC2 actions.
  - Define rules that match changes in your AWS environment and route these events to targets for processing.

Section 3: Amazon EC2 Auto Scaling

aws academy

Why is scaling important?

Used capacity (demand) Provisioned capacity

Amazon EC2 Auto Scaling

Used capacity (demand)

Provisioned capacity

- Helps you maintain application availability

- Enables you to automatically add or remove EC2 instances according to conditions that you define

- Detects impaired EC2 instances and unhealthy applications, and replaces the instances without your intervention

- Provides several scaling options – Manual, scheduled, dynamic or on-demand, and predictive

Typical weekly traffic at Amazon.com

Provisioned capacity

Sunday   Monday   Tuesday   Wednesday   Thursday   Friday   Saturday

November traffic to Amazon.com
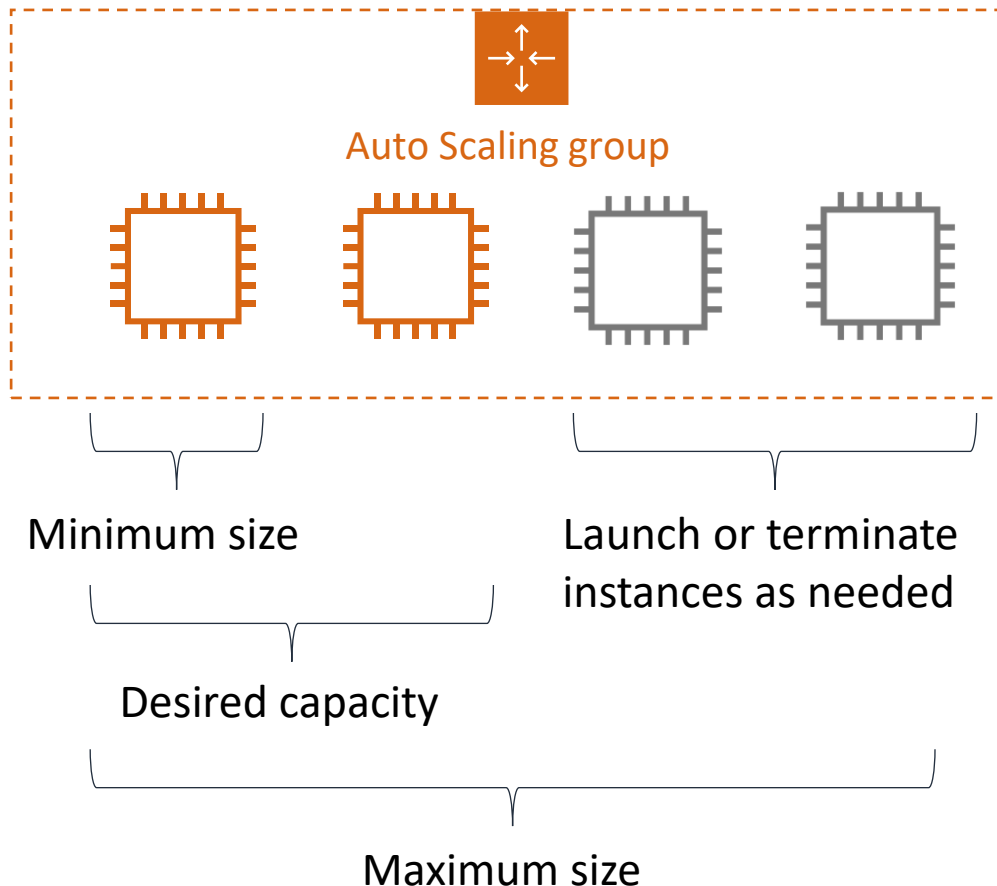
Provisioned capacity

76 percent

The challenge is to efficiently guess the unknown quantity of how much compute capacity you need.

24 percent
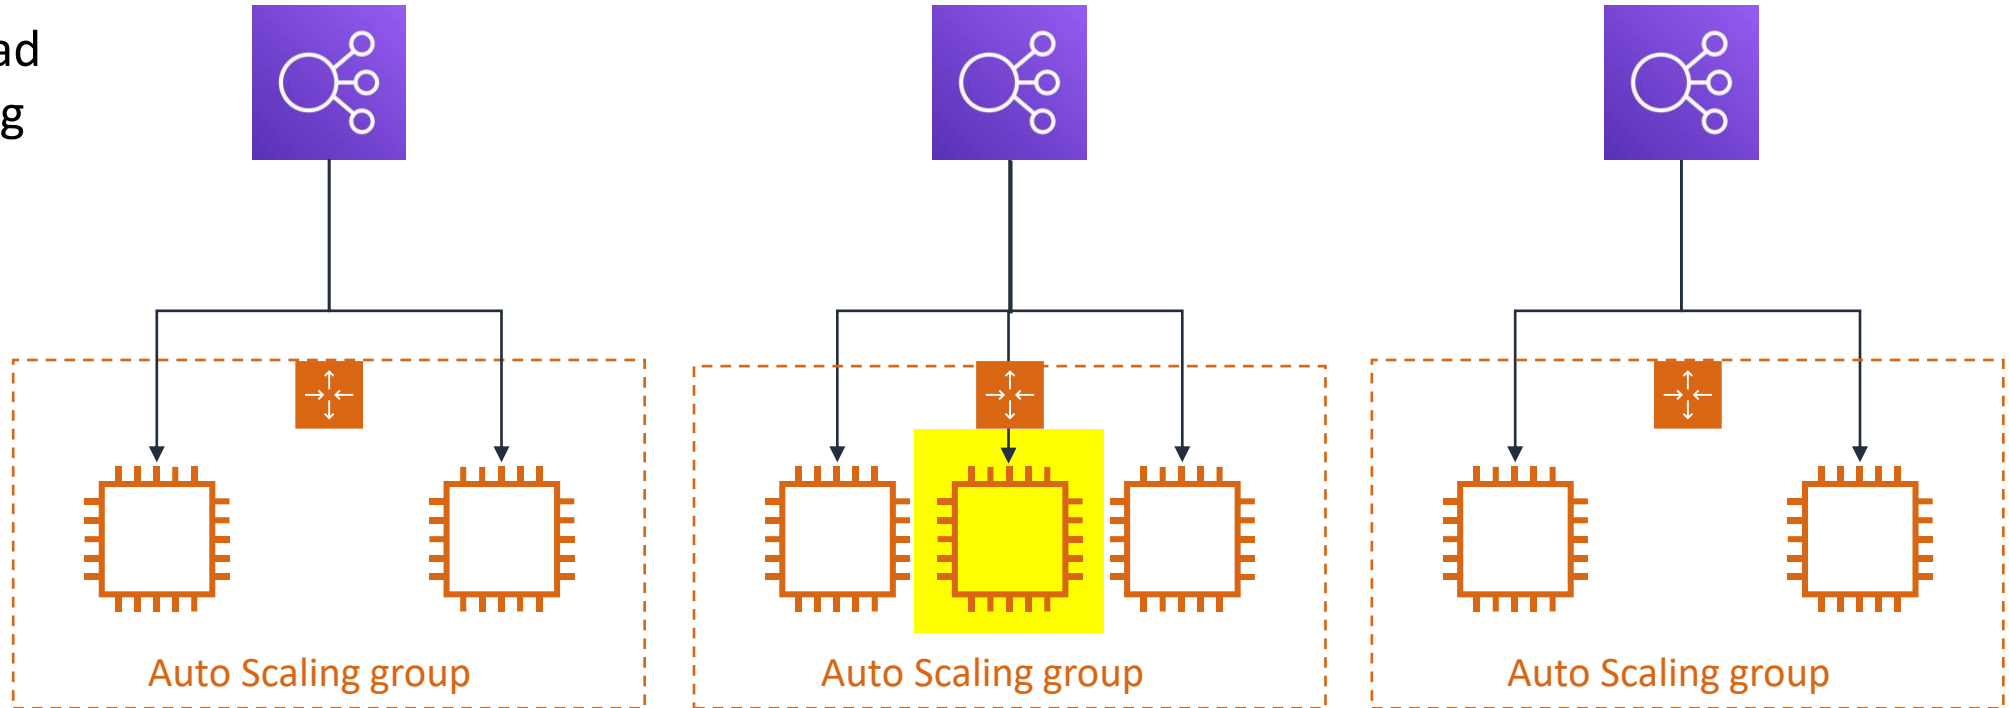
November

Auto Scaling groups

Auto Scaling group

Minimum size

Desired capacity

Maximum size

Launch or terminate instances as needed

An **Auto Scaling group** is a collection of EC2 instances that are treated as a logical grouping for the purposes of automatic scaling and management.

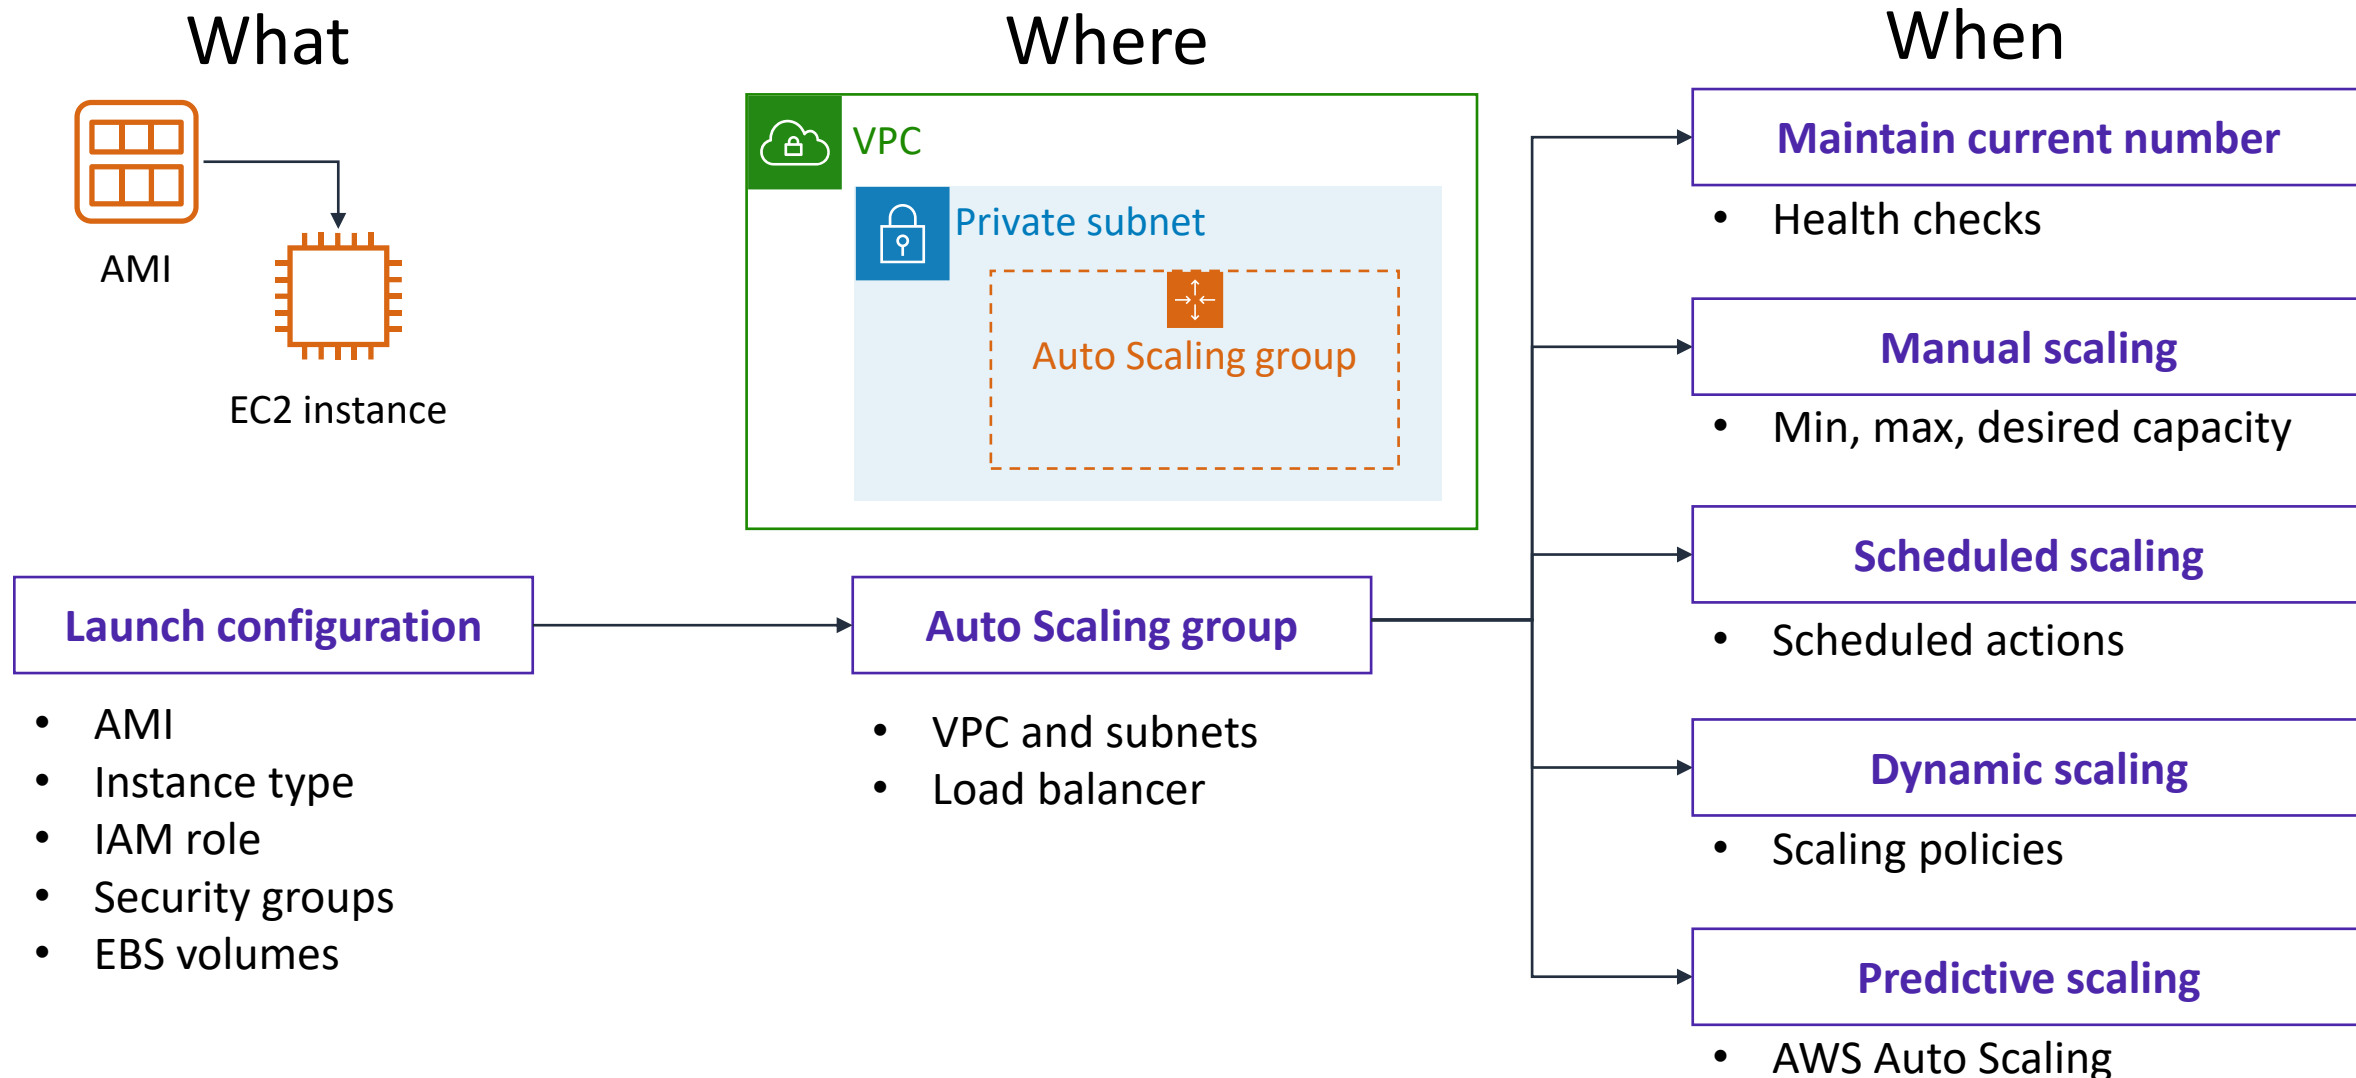# Scaling out versus scaling in

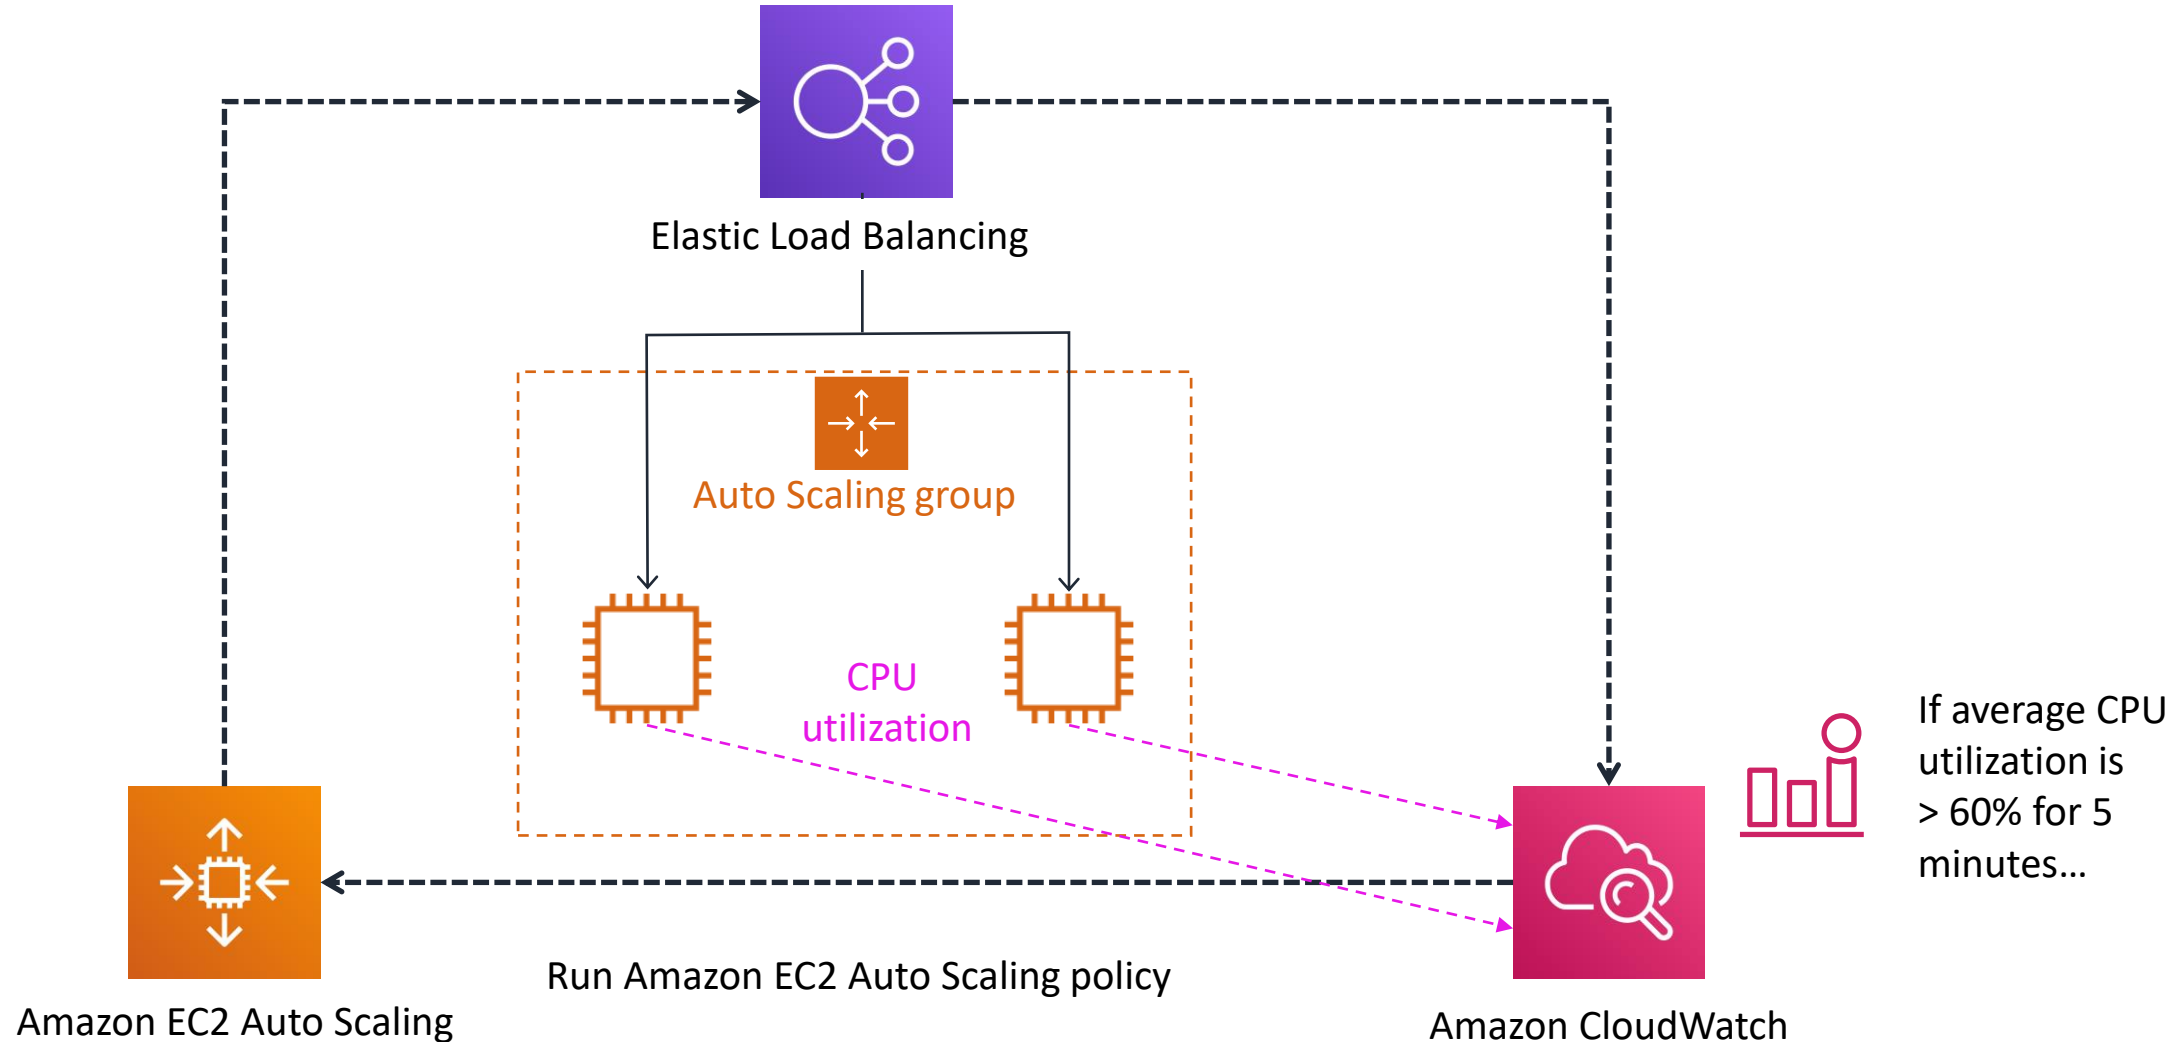Elastic Load Balancing

Base configuration

Scale out
(launch instances)

Scale in
(terminate instances)

Auto Scaling group

Auto Scaling group

Auto Scaling group

How Amazon EC2 Auto Scaling works

What

AMI

EC2 instance

Launch configuration

- AMI
- Instance type
- IAM role
- Security groups
- EBS volumes

Where

VPC

Private subnet

Auto Scaling group

Auto Scaling group

- VPC and subnets
- Load balancer

When

Maintain current number

- Health checks

Manual scaling

- Min, max, desired capacity

Scheduled scaling

- Scheduled actions

Dynamic scaling

- Scaling policies

Predictive scaling

- AWS Auto Scaling

Implementing dynamic scaling

Elastic Load Balancing

Auto Scaling group

CPU utilization

Amazon EC2 Auto Scaling

Run Amazon EC2 Auto Scaling policy

Amazon CloudWatch

If average CPU utilization is > 60% for 5 minutes…

AWS Auto Scaling

AWS Auto Scaling

- Monitors your applications and automatically adjusts capacity to maintain steady, predictable performance at the lowest possible cost

- Provides a simple, powerful user interface that enables you to build scaling plans for resources, including –
  - Amazon EC2 instances and Spot Fleets
  - Amazon Elastic Container Service (Amazon ECS) Tasks
  - Amazon DynamoDB tables and indexes
  - Amazon Aurora Replicas

Section 3 key takeaways

- Scaling enables you to respond quickly to changes in resource needs.

- Amazon EC2 Auto Scaling maintains application availability by automatically adding or removing EC2 instances.

- An Auto Scaling group is a collection of EC2 instances.

- A launch configuration is an instance configuration template.

- Dynamic scaling uses Amazon EC2 Auto Scaling, CloudWatch, and Elastic Load Balancing.

- AWS Auto Scaling is a separate service from Amazon EC2 Auto Scaling.

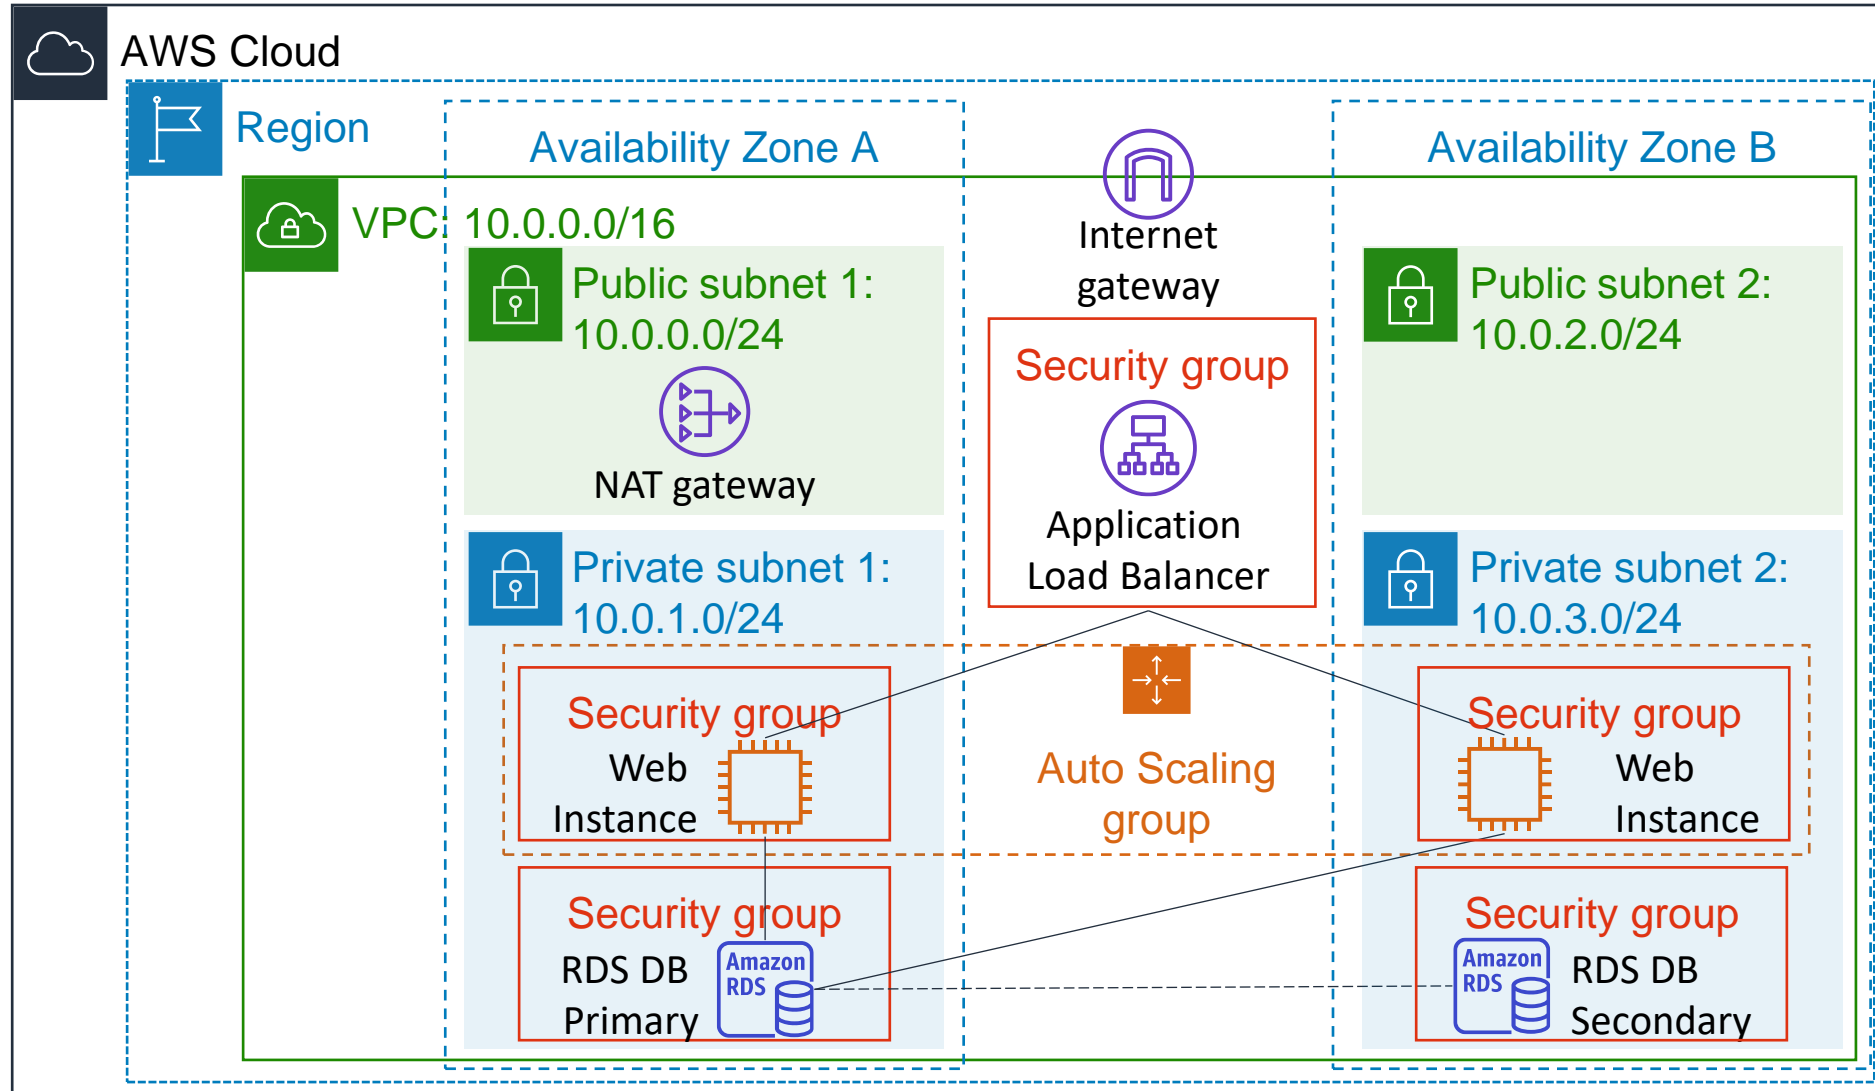# Lab 6:
Scale and Load Balance Your Architecture

Lab 6: Scenario

Lab 6: Tasks
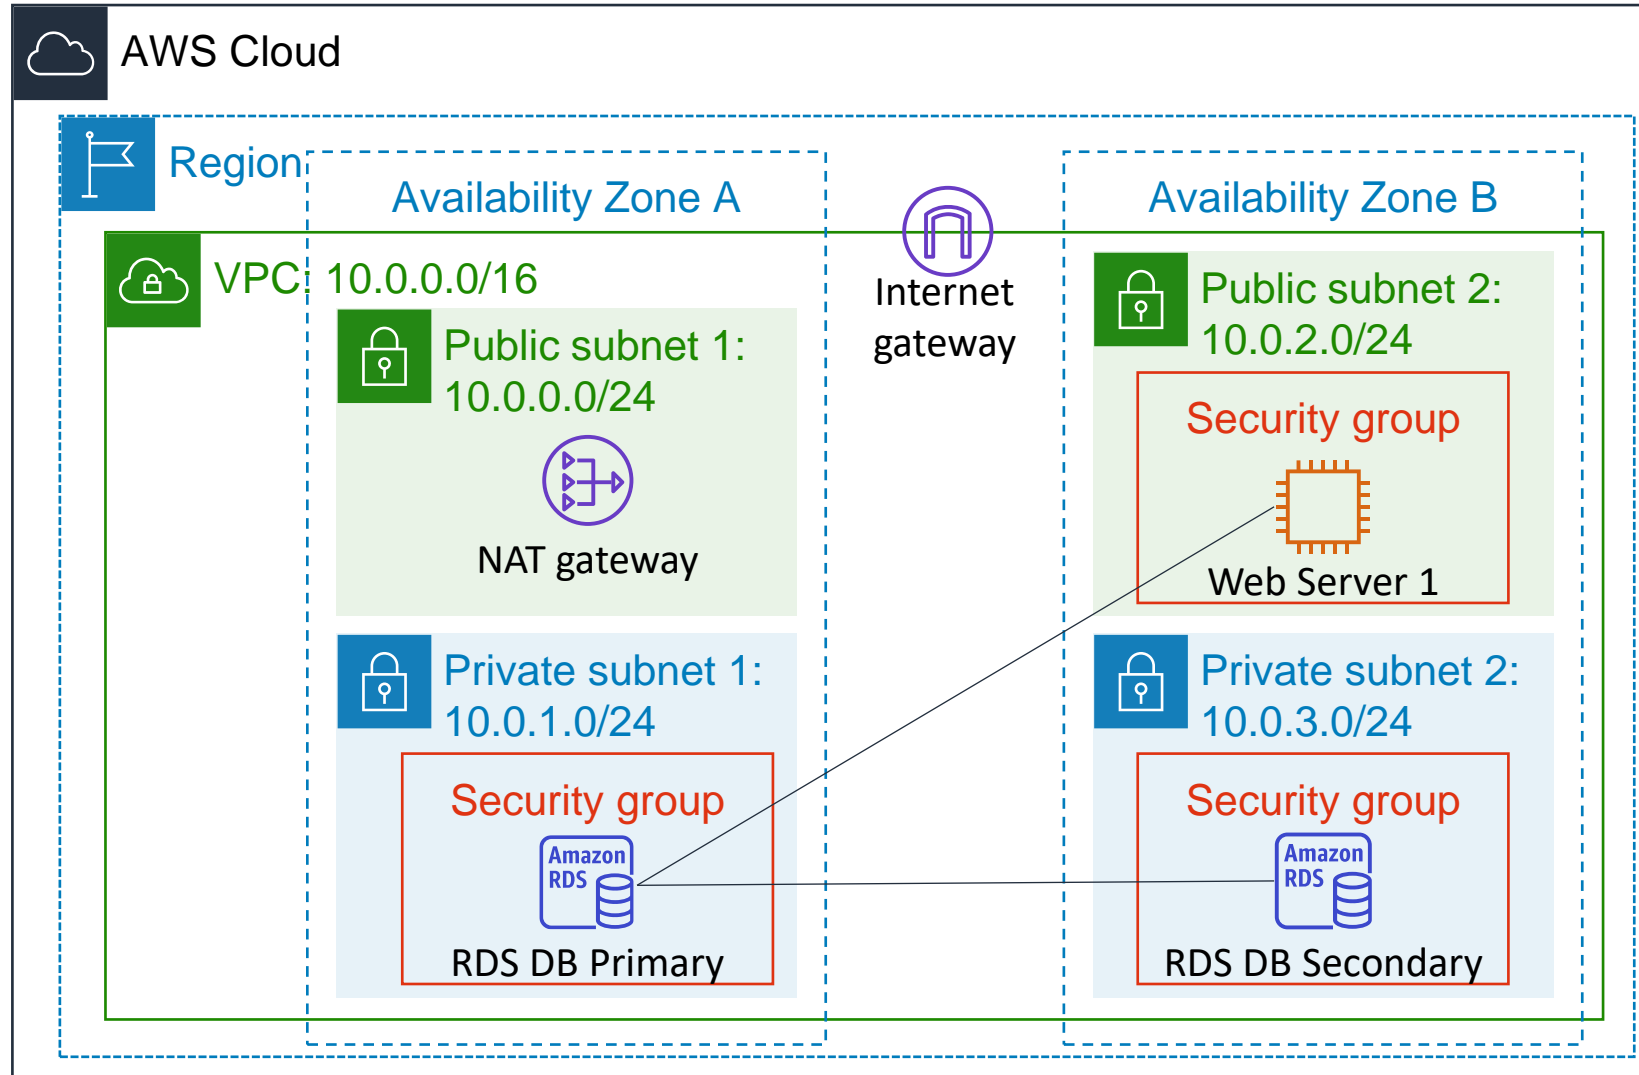
- Create an Amazon Machine Image (AMI) from a running instance.

- Create an Application Load Balancer.

- Create a launch configuration and an Auto Scaling group.

- Automatically scale new instances within a private subnet.

- Create Amazon CloudWatch alarms and monitor performance of your infrastructure.

Lab 6: Final product

~ 30 minutes

Begin Lab 6: Scale and Load Balance Your Architecture

Lab debrief:
Key takeaways

Module 10: Automatic Scaling and Monitoring

Module wrap-up

aws academy

Module summary

In summary, in this module you learned how to:

- Indicate how to distribute traffic across Amazon Elastic Compute Cloud (Amazon EC2) instances using Elastic Load Balancing.

- Identify how Amazon CloudWatch enables you to monitor AWS resources and applications in real time.

- Explain how Amazon EC2 Auto Scaling launches and releases servers in response to workload changes.

- Perform scaling and load balancing tasks to improve an architecture.

Complete the knowledge check

Sample exam question

Which service would you use to send alerts based on Amazon CloudWatch alarms?

A. Amazon Simple Notification Service

B. AWS CloudTrail

C. AWS Trusted Advisor

D. Amazon Route 53

Thank you

aws academy